

Portland State University
ECE 587/687

Caches and Memory-Level Parallelism

© Copyright by Alaa Alameldeen and Haitham Akkary 2015

Revisiting Processor Performance

- Program Execution Time =
(CPU clock cycles + Memory stall cycles)
x clock cycle time
- For each instruction:
CPI = CPI(Perfect Memory)
+ Memory stall cycles per instruction
- With no caches, all memory requests require accessing main memory
 - Very long latency (discuss)
- Caches filter out many memory accesses
 - Reduce execution time
 - Reduce memory bandwidth

Portland State University – ECE 587/687 – Spring 20152

Cache Performance

- Memory stall cycles Per Instruction =
Cache Misses per instruction x miss penalty
- Processor Performance:
CPI = CPI(Perfect Cache)
+ miss rate x miss penalty
- Average memory access time =
Hit ratio x Hit latency + Miss ratio x Miss penalty
- Cache hierarchies attempt to reduce average memory access time

Portland State University – ECE 587/687 – Spring 20153

Cache Performance Metrics

- Hit ratio: #hits / #accesses
- Miss ratio: #misses / #accesses
- Miss rate: Misses per instruction (or 1000 instructions)
 - Miss rate = miss ratio x memory accesses per inst
- Hit time: time from request issued to cache until data is returned to the processor
 - Depends on cache design parameters
 - Bigger caches, larger associativity, or more ports increase hit time
- Miss penalty: depends on memory hierarchy parameters

Portland State University – ECE 587/687 – Spring 20154

Why Do Caches Work?

- Spatial Locality
 - If data at a certain address is accessed, it is likely that data located at nearby addresses will also be accessed in the (near) future
 - Implication: Cache line (block) size tradeoff
- Temporal Locality
 - If data at a certain address is accessed, it is likely the same data will be accessed in the (near) future
 - Implication: Replacement algorithms try to predict which lines will be accessed

Portland State University – ECE 587/687 – Spring 20155

Memory Hierarchy

- First-level caches
 - Usually Split I & D caches
 - Small and fast
- Second-level caches
 - Usually on-die
 - SRAM cells
- Third-level... etc.?
- Main memory
 - DRAM cells
 - focus on density
- Solid-State Disk?
- Hard Disk
 - Usually magnetic device, non-volatile
 - Slow access time

```
graph TD; Processor[Processor] <--> L1I[L1I$]; Processor <--> L1D[L1D$]; L1I <--> L2[L2 Cache]; L1D <--> L2; L2 <--> L3[L3/LLC]; L3 <--> MainMemory[Main Memory]; MainMemory <--> Disk[(Disk)];
```

Portland State University – ECE 587/687 – Spring 20156

Basic Cache Structure

- Array of blocks (lines)
 - ◆ Each block is usually 32-128 bytes
- Finding a block in cache:

Data Address	Tag	Index	Offset
--------------	-----	-------	--------
- Offset: byte offset in block
- Index: Which set in the cache is the block located
- Tag: Needs to match address tag in cache

Portland State University – ECE 587/687 – Spring 20157

Associativity

- Set associativity
 - ◆ Set: Group of blocks corresponding to same index
 - ◆ Each block in the set is called a Way
 - ◆ 2-way set associative cache: each set contains two blocks
 - ◆ Direct-mapped cache: each set contains one block
 - ◆ Fully-associative cache: the whole cache is one set
- Need to check all tags in a set to determine hit/miss status

Portland State University – ECE 587/687 – Spring 20158

Example: Cache Block Placement

- Consider a 4-way, 32KB cache with 64-byte lines
- Where is 48-bit address 0x0000FFFFAB64?
 - ◆ Number of lines = cache size / line size = 32K / 64 = 512
 - ◆ Each set contains 4 lines \Rightarrow Number of sets = $512/4 = 128$ sets
 - ◆ Offset bits = $\log_2(64) = 6$: 0x24
 - ◆ Index bits = $\log_2(128) = 7$: 0x2D
 - ◆ Tag bits = $48-(6+7) = 35$: 0x00007FFD

Portland State University – ECE 587/687 – Spring 20159

Types of Cache Misses

- Compulsory (cold) misses: First access to a block
 - ◆ Prefetching can reduce these misses
- Capacity misses: A cache cannot contain all blocks needed in a program - some blocks are discarded then accessed
 - ◆ Replacement policies should target blocks that won't be used later
- Conflict misses: Blocks mapping to the same set may be discarded (in direct-mapped and set-associative caches)
 - ◆ Increasing associativity can reduce these misses
- For multiprocessors, coherence misses can also happen
 - ◆ Details in ECE 588/688

Portland State University – ECE 587/687 – Spring 201510

Cache Replacement and Insertion Policies

- Cache replacement policy:
 - ◆ On a cache line fill, which victim line to replace
 - ◆ Only applicable to set-associative caches
 - ◆ Examples: LRU, more advanced policies
 - ◆ Discuss stack algorithms
- Cache insertion policy:
 - ◆ When a cache line is filled, what would be its priority in the replacement stack
 - ◆ LRU: fill line is inserted in "Most Recently Used" position
 - ◆ Other policies: LIP, BIP, DIP
 - ◆ Dead block prediction helps determine lines that won't be reused

Portland State University – ECE 587/687 – Spring 201511

Non-Blocking Cache Hierarchy

- Superscalar processors require parallel execution units
 - ◆ Multiple pipelined functional units
 - ◆ Cache hierarchies capable of simultaneously servicing multiple memory requests
 - Do not block cache references that do not need the miss data
 - Service multiple miss requests to memory concurrently
 - ◆ Revisit miss penalty with memory-level parallelism

Portland State University – ECE 587/687 – Spring 201512

Miss Status Holding (Handling) Registers

- MSHRs facilitate non-blocking memory level parallelism
- Used to track address, data, and status for multiple outstanding cache misses
- Need to provide correct memory ordering, respond to CPU requests, and maintain cache coherence
- Design details vary widely between different processors
 - ◆ But basic functions are similar

Cache & MSHR Organization

- Paper Fig 1: block diagram of cache organization
- Main Components:
 - ◆ MSHR: One register for each miss to be handled concurrently
 - ◆ N-way comparator: Compares an address to all block addresses in MSHRs ($N = \#MSHRs$)
 - ◆ Input Stack: Buffer space for all misses corresponding to MSHR entries
 - Size = $\#MSHRs \times \text{block size}$
 - ◆ Status update and collecting networks
- Current implementations combine the MSHR and input stack

MSHR Structure

- Each MSHR contains the following information
 - ◆ Data address
 - ◆ PC of requesting instruction
 - ◆ Input identification tags & Send-to-CPU flags: Allows forwarding requested word(s) to the CPU
 - ◆ Input stack indicators (flags): Allows reading data directly from input stack
 - ◆ Partial write codes: Indicates which bytes in a word has been written to the cache
 - ◆ Valid flag
 - ◆ Obsolete flag: either info is not valid or MSHR hit on data in transit

MSHR Operation

- On a cache miss, one MSHR is assigned
 - ◆ Valid flag set
 - ◆ Obsolete flag cleared
 - ◆ Data address saved
 - ◆ PC of requesting instruction saved
 - ◆ Appropriate send-to-CPU flags set and others cleared
 - ◆ Input identification tag saved in correct position
 - ◆ Partial write codes cleared
 - ◆ All MSHRs pointing to the same data address purged
- Optimal number of MSHRs: paper figure 2

Virtual vs. Physical Addressing

- Using virtual addresses to access the L1 cache reduces latency
 - ◆ Physical addresses need address translation
- However, virtually-addresses caches introduce problems
 - ◆ Handling synonyms: multiple VAs mapping to same PA
 - ◆ Address translation needed on L1 misses
 - ◆ Reverse translation needed for coherence in a multiprocessor system
 - ◆ Need to invalidate whole cache on a context switch
- Two-level hierarchies can help achieve the best of both worlds (Wang paper Figure 1)

Reducing Cache Misses

- Cache misses are very costly
- Need multiple cache levels with
 - ◆ High associativity or/and victim caches to reduce conflict misses
 - ◆ Effective replacement algorithms
 - ◆ Data and instruction prefetch
 - Preferably with multiple stream buffers
- More details in Wednesday's class

Reading Assignment

- N. Jouppi, "Improving Direct-Mapped Cache Performance by the Addition of a Small Fully-Associative Cache and Prefetch Buffers," ISCA 1990 (Review)
- T.F. Chen and J.L Baer, "Effective Hardware-Based Data Prefetching for High-Performance Processors," IEEE Transactions on Computers, 1995 (Skim)